

Energy-based models (EBM), part 2

Aapo Hyvärinen

for the NNDL course, 16 April 2025 (first half)

- ▶ Noise-contrastive estimation
- ▶ Comparison of estimators

Literature: Murphy's book Chapter 24

From last time: Estimation of energy-based models (EBM)

- ▶ Also called estimation of *unnormalized* statistical models
- ▶ Density function is known only up to a multiplicative constant

$$p_{\text{norm}}(\mathbf{x}; \boldsymbol{\theta}) = \frac{1}{Z(\boldsymbol{\theta})} p_{\text{un}}(\mathbf{x}; \boldsymbol{\theta})$$

$$Z(\boldsymbol{\theta}) = \int_{\boldsymbol{\xi} \in \mathbb{R}^n} p_{\text{un}}(\boldsymbol{\xi}; \boldsymbol{\theta}) d\boldsymbol{\xi} \quad (1)$$


with p_{norm} : actual density; p_{un} : unnormalized version

- ▶ Functional form of p_{un} is “known”
(i.e. can be easily computed)
- ▶ Normalization constant (= “partition function”) Z
“unknown”, i.e. *cannot be easily computed*
- ▶ Computation of Z based on the above would need numerical integration: very difficult in high dimensions


Noise-contrastive estimation (NCE) as already seen in SSL introduction

- ▶ Train a nonlinear classifier to discriminate observed data from some artificial noise
- ▶ To be successful, the classifier must “discover structure” in the data
- ▶ For example, compare natural images with Gaussian noise

Natural images



Gaussian noise



- ▶ A nice intuitive idea, but what does this system actually do?
 - ▶ ... in addition to learning nice features for supervised learning.

Definition of classifier in NCE

- ▶ Observed data set $\mathbf{X} = (\mathbf{x}(1), \dots, \mathbf{x}(N))$ has *unknown* pdf $p_{\mathbf{x}}$
- ▶ Generate “noise” $\mathbf{Y} = (\mathbf{y}(1), \dots, \mathbf{y}(N))$ with known pdf $p_{\mathbf{y}}$
- ▶ Define a nonlinear function, a NN, $g(\mathbf{u}; \theta)$, which models data log-density $\log p_{\mathbf{x}}(\mathbf{u})$.
 - ▶ The function g corresponds to $\log p_{\text{un}}$
 - ▶ here we initially focus on general density estimation
- ▶ We use standard logistic regression, but with the special nonlinear function

$$G(\mathbf{z}; \theta) = g(\mathbf{z}; \theta) - \log p_{\mathbf{y}}(\mathbf{z}). \quad (2)$$

where \mathbf{z} takes values over the whole data (either \mathbf{x} or \mathbf{y}).

- ▶ Next we prove that such training solves EBM estimation: θ converges to the true values that generated the data, or: in case of pdf approximation, g converges to log-pdf of \mathbf{x} .

Definition of logistic regression and its objective

- ▶ Begin by transforming probability p to $p/(1 - p)$, “odds ratio”
- ▶ Logistic regression uses function r to model *log-odds*

$$r(\mathbf{u}; \boldsymbol{\theta}) \approx \log \frac{p(c = 1 | \mathbf{z} = \mathbf{u})}{1 - p(c = 1 | \mathbf{z} = \mathbf{u})} \quad (3)$$

where $c \in \{1, 2\}$ is class label.

- ▶ To derive likelihood, first solve $p(c = 1 | \mathbf{z} = \mathbf{u})$ as

$$p(c = 1 | \mathbf{z} = \mathbf{u}) = \sigma(r(\mathbf{u}; \boldsymbol{\theta})) \quad (4)$$

whose solution uses the “logistic function” $\sigma(r) = \frac{1}{1 + \exp(-r)}$.

- ▶ Likewise we solve: $p(c = 2 | \mathbf{z} = \mathbf{u}) = 1 - \sigma(r(\mathbf{u}; \boldsymbol{\theta}))$.
- ▶ Take sum of logarithms of the likelihood for the two classes
- ▶ This gives the standard log-likelihood objective for training:

$$J(\boldsymbol{\theta}) = \sum_{i=1}^N \log [\sigma(r(\mathbf{x}(i); \boldsymbol{\theta}))] + \log [1 - \sigma(r(\mathbf{y}(i); \boldsymbol{\theta}))] \quad (5)$$

where \mathbf{x} is data in class 1 and \mathbf{y} in class 2; classes of equal size

What does logistic regression actually do?

- Fundamental *Theorem*: logistic regression trained with Eq. (5) learns to approximate difference of log-densities:

$$r(\mathbf{u}; \boldsymbol{\theta}) \rightarrow \log p_{\mathbf{x}}(\mathbf{u}) - \log p_{\mathbf{y}}(\mathbf{u}) \quad (6)$$

which becomes exact with infinite sample size and universal approximation capability of function approximator r

- Heuristic proof: re-interpret the log-odds as *log density ratio*

$$\log \frac{p(c = 1 | \mathbf{z} = \mathbf{u})}{1 - p(c = 1 | \mathbf{z} = \mathbf{u})} = \log p(c = 1 | \mathbf{z} = \mathbf{u}) - \log p(c = 2 | \mathbf{z} = \mathbf{u})$$

since obviously $p(c = 2 | \mathbf{z} = \mathbf{u}) = 1 - p(c = 1 | \mathbf{z} = \mathbf{u})$.

- This is what r tries to model to according to definition
- We can manipulate this by using
$$p(c = C | \mathbf{z} = \mathbf{u}) = p(\mathbf{z} = \mathbf{u} | c = C) p(c = C) / p(\mathbf{z} = \mathbf{u})$$
- $p(c = C)$ and $p(\mathbf{z} = \mathbf{u})$ are equal for $C \in \{1, 2\}$ and cancel
- The above equals $\log p(\mathbf{z} = \mathbf{u} | c = 1) - \log p(\mathbf{z} = \mathbf{u} | c = 2)$
- $\log p_{\mathbf{x}}(\mathbf{u}) - \log p_{\mathbf{y}}(\mathbf{u})$ is the same in different notation

So, what does the classifying system do in NCE?

- ▶ Recall: NCE is logistic regression with regression function $r(\mathbf{u}; \boldsymbol{\theta}) := G(\mathbf{u}; \boldsymbol{\theta}) := g(\mathbf{u}; \boldsymbol{\theta}) - \log p_{\mathbf{y}}(\mathbf{u})$
- ▶ Fundamental NCE Theorem:
 - ▶ Assume $g(\mathbf{u}; \boldsymbol{\theta})$, i.e. the NN trained in NCE, can approximate any function, and the sample size is infinite
 - ▶ Then, the maximum of classification objective is attained when

$$g(\mathbf{u}; \boldsymbol{\theta}) = \log p_{\mathbf{x}}(\mathbf{u}) \quad (7)$$

where $p_{\mathbf{x}}(\mathbf{u})$ is the pdf of the observed data.

- ▶ Proof: We just saw that logistic regression learns to approximate difference of log-densities (here: data vs. noise), in the limit:

$$G(\mathbf{u}; \boldsymbol{\theta}) = \log p_{\mathbf{x}}(\mathbf{u}) - \log p_{\mathbf{y}}(\mathbf{u}) \quad (8)$$

So, comparing with the definition of G , we have Eq. (7).

Connection to EBM: NCE learns normalization by itself

- ▶ The maximum of objective function is attained when $g(\mathbf{u}; \theta) = \log p_{\mathbf{x}}(\mathbf{u})$,
and there is *no constraint* on g in this optimization problem!
 - ▶ In particular, no normalization constraint
(such as $\int \exp(g(\mathbf{u}; \theta)) d\mathbf{u} = 1$)
- ▶ Even if the family $g(\mathbf{u}; \theta)$ is not normalized (and a NN is usually not), the maximum is still attained for the properly normalized pdf
- ▶ In practice, normalization constant (partition function) will be estimated together with the other parameters
- ▶ Such a new normalization parameter is learned in a NN with universal approximation
- ▶ In non-NN case where $g(\mathbf{x}; \theta)$ is a model with a limited number of parameters:
 - ▶ Add a new parameter c and use $g(\mathbf{u}; \theta) + c$
 - ▶ Corollary to NCE theorem: If data generated according to model, i.e. $\log p_{\mathbf{x}}(\mathbf{u}) = g(\mathbf{u}; \theta^*)$, this gives a *statistically consistent* estimator for θ^* in unnormalized case.

Estimating parameters vs. estimating density

- ▶ Classical estimation theory considers finding *parameters* θ
 - ▶ That was the approach in score matching slides
- ▶ But with NCE, we focused on estimating the *density*
- ▶ In statistics literature, these are two very different problems: Density estimation is called “non-parametric”, actually meaning that the number of parameters is infinite (function spaces such as L^2 have an infinite dimension)
- ▶ When you learn a NN, clearly the number of parameters is finite, but often huge (practically infinite?)
- ▶ In EBM, we can interchangeably talk about estimation of parameters or density, the methods are the same:
 - ▶ Using a big NN as $g \approx$ non-parametric estimation (of density)
 - ▶ Using a carefully chosen function as $g \approx$ parametric estimation

Choice of noise distribution in NCE

- ▶ The noise distribution p_y is an important design parameter.
- ▶ We would like to have p_y which fullfills the following:
 1. Easy to sample from
 - ▶ But we only need to sample noise once, off-line
 2. Has an analytical properly normalized expression
 - ▶ But we only need to, e.g., normalize it once
 3. It leads to (e.g.) a small mean-squared error of the estimator
 - ▶ This can be analyzed, but not simple to find good distribution
- ▶ Intuitively, noise should be rather similar to data, so that classification not too easy
- ▶ In practice, we can take Gaussian noise with the same mean and covariance as the data.
 - ▶ Far from optimal, but simple

Comparison between different estimators

- ▶ Given two estimators for the same model (e.g. SM and NCE), how can we say one of them is better?
 - ▶ We have two kinds of properties to compare
 - ▶ Statistical performance
 - ▶ Computation
 - ▶ Statistics: Estimation theory compares estimators e.g. by
 1. Consistency: Does the estimator converge to the true θ^* with infinite sample size?
 2. Mean squared error (MSE): $E\{\|\hat{\theta} - \theta^*\|^2\}$
where expectation is taken over sample size which is fixed.
 - ▶ We prefer a consistent estimator over an inconsistent one
 - ▶ We prefer estimator with smaller MSE
 - ▶ Computation: Typically only computation time considered
 - ▶ Could consider energy consumption, memory requirements, etc.
 - ▶ We see that there are many properties to compare
- Usually, some kind of trade-off, arbitrary choices necessary

General comparison between score matching and NCE

- ▶ Consider a model $p(\mathbf{x}; \boldsymbol{\theta})$, and its estimation by SM or NCE

Statistics

- ▶ Both NCE and basic SM are consistent
- ▶ Denoising SM “consistently” estimates noisy pdf, not original
- ▶ MSE of NCE depends heavily on choice of noise distribution
- ▶ (In denoising SM, must choose noise variance but this is usually easy to optimize by trying out different values.)

Computation

- ▶ In a NN, basic SM algebraically/computationally difficult
- ▶ Denoising score matching works well (empirically) with NN's
- ▶ Not much difference between NCE and DSM

... How about maximum likelihood? Statistically the best, but computationally worst for EBM (needs numerical integration)

Validation and comparison of statistical models

- ▶ Above, we compared different estimation principles for a single model
- ▶ We might also want to compare two models using same estimation principle (and same data)
 - Possible by comparing the values of the objective function for the two models (e.g. better fit to data score function)
 - ▶ Importantly, this should be on separate (held-out) test data
- ▶ One problem: we don't really know the baseline
 - ▶ What kind of value for objective is good at all? Difficult to say.
 - ▶ This is a problem even with MLE:
no clue what value of likelihood is “good” (“enough”)
 - ▶ So, difficult to say when a statistical model is good or bad
 - ▶ In contrast to classification with two classes: it is intuitively clear that 52% accuracy is not good, while 95% accuracy is very good

Different ingredients in probabilistic modelling

- ▶ In early lectures, DL was characterized as
 1. Objective function J , based on learning principle
 2. NN as function approximator (\mathbf{g}) with specified architecture
 3. Optimization algorithm applied on J
- ▶ We have now seen how the construction of objective function happens in EBM. We choose:
 - A Estimation principle e.g. MLE, score matching, NCE, etc.
 - B Architecture, which equivalently defines *Statistical model*
 - ▶ since we define e.g. $\log p(\mathbf{x}; \theta) = \mathbf{g}(\mathbf{x}; \theta)$ in DL case
- ▶ These two together give the *Objective function* J
- ▶ We still need
 - C Optimization algorithm applied on J
- ▶ This logic applies also to latent variable models (next lecture)

Conclusion (of EBM parts 1 and 2)

- ▶ We saw three methods for estimating parameters in unnormalized (energy-based) models; MLE being too difficult
- ▶ These methods avoided the very expensive numerical integration needed to compute the normalization constant
- ▶ In score matching, match gradients of log-densities
 - ▶ normalization constant is completely avoided
- ▶ In denoising score matching, train autoencoder to reconstruct data from noisy samples
 - ▶ computationally simpler than basic SM in an NN
- ▶ In noise-contrastive estimation, learn logistic regression to discriminate data from artificial noise
 - ▶ normalization constant can be estimated like any parameter
- ▶ Increasingly used in deep learning as general density approximators
- ▶ Basis of state-of-the-art image generation methods (see next lecture)